

# Computing in C for Science

Lecture 1 of 5

Dr. Steven Capper

<http://www2.imperial.ac.uk/~sdc99/ccourse/>  
steven.capper99@imperial.ac.uk

16<sup>th</sup> November 2011

Imperial College  
London

Steven Capper

C for Science - Lecture 1

## Introduction to the Course

This course is based on the C course written by my PhD supervisor Dan Moore: <http://www.ma.ic.ac.uk/~drmi>

### Aims of the Course

- To introduce modern C programming from scratch and,
- provide insight into scientific computing (floating point arithmetic, optimisation, ...).

### Five lectures spread over five weeks.

- Each lecture will take  $\approx$  1 hour,
- and involve at least an hour of practical work.
- This is very intense.
- Please feel free to ask questions outside course hours.

Steven Capper

C for Science - Lecture 1

## A Rough History of C

### Invented $\approx$ 1970

By Dennis Ritchie working in Bell Labs USA; to facilitate development of a portable UNIX.

### C has been standardised

- 1989 ANSI standard ratified *ANS X3.159-1989*.
- 1990 ISO standard *ISO/IEC 9899:1990*. Aka *C90*.
- 2000 ISO standard *ISO/IEC 9899:1999*. Aka *C99*.

### C has evolved into C++

Bjarne Stroustrup developed C++ (C with class). Unlike C, C++ is still under very active development (C++11 being the most recent standard at the time of writing).

Steven Capper

C for Science - Lecture 1

## What are C and C++?

- C is a cross-platform, compiled, general-purpose language.
- C++ can loosely be thought of as C's object oriented big brother.

The vast majority of the programs running on your computer (including the operating system kernel), are written in either C or C++.

Steven Capper

C for Science - Lecture 1

## Why Use C? (Over Maple, Matlab, S-Plus...)

### Speed

C programs are compiled to machine code, the resulting routines *can* run several orders of magnitude quicker than their equivalents in interpreted environments.

### Flexibility

The C language is intrinsically low level, one can manipulate complex data structures with surprisingly little code.

### Portability

A well written C program can target many different environments (Windows PCs, Linux workstations, Apple Macs, DEC Alphas, Embedded devices, ...).

## Getting Started

You will need:

- A C compiler (many different ones to choose from, some are free).
- Some documentation (such as the lecture notes/exercises from this course, a good book, online guides).
- Lots, and lots of time.

## Free C compilers

Free as in free for academic/commercial use

### Linux/UNIX

- gcc - The GNU Compiler Collection, C compiler.  
<http://gcc.gnu.org>.

### Windows

- cygwin - A set of GNU libraries ported to Windows (free usage restricted to GPL apps),  
<http://www.cygwin.com/>.
- MinGW - Minamalist GNU for Windows (no restrictions),  
<http://www.mingw.org/>.
- Visual C++ 2008 Express Edition - Microsoft's free compiler (no restrictions),  
<http://www.microsoft.com/express/vc/>

## Some Free IDEs

gcc is a command line driven program. Thankfully, there exist many free Integrated Development Environments (IDEs) that simplify the development process. Popular IDEs include:

- Eclipse - For Windows/UNIX. Primarily used for Java, but is good for C development too.  
<http://www.eclipse.org/>
- NetBeans - For Windows/UNIX. Another Java IDE with C development functionality. <http://netbeans.org/>
- Xcode Tools - For Apple Macs. The development environment used by Apple.  
<http://developer.apple.com/technology>

## Some Non-Free C Compilers

- Borland/Inprise/Borland/Code Gear C++ Builder/Turbo C++ - Can be found at: <http://www.codegear.com>
- Intel - for Windows/Linux (will compile good code for AMD processors too!). Free for personal use, academic/commercial licenses obtainable from: <http://www.polyhedron.com>
- SilverFrost(Salford) - for Windows, includes a C compiler, personal evaluation version from: [http://www.silverfrost.com/32/ftn95/personal\\_edition/](http://www.silverfrost.com/32/ftn95/personal_edition/)
- Microsoft Visual Studio 2008 Professional - Microsoft's flagship compiler. Ninety day free trial available: <http://msdn2.microsoft.com/en-us/vstudio/product>

## Calling all Students!

### Microsoft DreamSpark

Microsoft have released the full Visual Studio 2008 Professional Edition and Server 2008 (and 2008 R2) for student use. <https://downloads.channel8.msdn.com/>

### Licensing

I have no idea how this is licensed, please do check before using it for research/commercial use!

## Books for C

### Kernighan and Ritchie (K&R2)



*The C Programming Language, Second Edition*, Prentice Hall. A very well-written, concise C reference.

### Numerical Recipes in C



By Press, Teukolsky, Vetterling & Flannery, **Second Edition**, CUP. Contains a lot of useful scientific computing information and provides high quality C example code. A free online edition can be found at: <http://www.nr.com>


## More Books for C/C++

### King



*C Programming A Modern Approach* Norton & Company. A well paced introduction to the C programming language with some introductory material to C++ too.

### Numerical Recipes in C++



By Press, Teukolsky, Vetterling & Flannery, **Third Edition**, CUP. A C++ only, heavily revised edition of the text. If you want to perform scientific computations in C++ rather than C, I would recommend this. More information at: <http://www.nr.com>

## Building a C Program

- To *build* an executable from source, we carry out the following three steps:

### Edit Source

Use a text editor to create a `.c` file.

### Compile

With a C compiler, this creates *object file(s)*.

### Link

Combine the object files together into an *executable*.

- These steps can be automated by *Integrated Development Environments* (IDEs).

## The Traditional Way to Start

```
1 #include <stdio.h>
2
3 int main(void)
4 {
5     printf("Hello World!\n");
6     return 0;
7 }
```

### The "Hello World" Program

A traditional first program started by Ritchie. This is one of the smallest possible C programs that demonstrates some functionality (printing to screen).

### Line 1

A *pre-processor directive* (it begins with a #) advertising extra routines to the compiler.

### Line 2

An empty line, or equivalently, a line consisting solely of *whitespace*. This is ignored by the compiler but makes the source code more readable.

### Line 3

A *function declaration*, defining our `main` function. The `main` function is where our program starts and is known as an *entry point*. Our `main` function takes *no parameters* (`void`) and *returns* an integer (`int`).

### Line 4

*Opening brace*, all statements enclosed between the braces `{`, `}` belong to the `main` function.

### Line 5

A *statement*; the `printf` (print formatted) function is called with the argument `"Hello World!\n"`. This prints:  
Hello World!  
to *standard output* (usually a text console).

### Line 6

A *return statement*, we exit `main` with a return code of 0. The system interprets 0 as "success".

### Line 7

A *closing brace*, everything after this line does not belong to `main`.

## Another C Program - What does this do?

```
1 #include <stdio.h>
2
3 int main(void)
4 {
5     int low=-40, high=140, step=5, f, c;
6     c = low;
7     while (c <= high)
8     {
9         f = 32+9*c/5;
10        printf("%6d \t %6d\n", c, f);
11        c = c + step;
12    }
13    return 0;
14 }
```

Steven Capper

C for Science - Lecture 1

### Lines 1, 2, 3 & 4

Identical meaning as in the previous program.

### Line 5

*Local variable declarations*; the integers `low`, `high`, `step`, `f` and `c` are declared. These are local to `main`. The variables `low`, `high` and `step` are *initialised* with the values; whilst `f` and `c` are *undefined*.

### Line 6

The local variable `c` is *assigned* the value of `low`.

### Lines 7, 8 & 12

A *while* loop is defined. For as long as the variable `c` is less than or equal to `high`, the code between the braces on lines 8 and 12 is executed.

Steven Capper

C for Science - Lecture 1

### Line 9

The local variable `f` is assigned a value from the integer arithmetic expression involving `c`.

### Line 10

The variables `c` and `f` are printed to standard out, each six characters wide, separated by a tab and two spaces.

### Line 11

The local variable `c` is incremented by `step`.

### Lines 13 & 14

Have an identical meaning as in the last program.

Steven Capper

C for Science - Lecture 1

## printf - declared in <stdio.h>

We call `printf` as follows:

```
printf(formatString, var1, var2, ..., varN);
```

where,

### formatString

The format string tells `printf` how many variables need printing. A format string can contain *format specifiers*, these tell `printf` exactly how to print out each variable, some examples:

`"%6d"` print out an integer (6 characters wide).

`"%g"` print out a floating point number.

### var1, ...

`printf` accepts a variable list of arguments, which can be of different type. Care must be taken to match `formatString` with the variables.

Steven Capper

C for Science - Lecture 1

## Special Characters

- The backslash `\` character in C has a special meaning, it is known as the *escape character*.
- We combine the escape character with other characters, to form an *escape sequence*, here are some examples:

<code>\n</code>	New line	<code>\f</code>	Form feed (new page)
<code>\t</code>	Tab	<code>\\</code>	<code>\</code>
<code>\b</code>	Backspace	<code>\"</code>	<code>"</code>
<code>\r</code>	Carriage return	<code>\'</code>	<code>'</code>
<code>\a</code>	Bell		

## Numbers in C - 2 General Types

- *Integers* - `short`, `unsigned short`, `int`, `unsigned int`, `long`, `unsigned long`, `long long`, `unsigned long long`. Integer types in C can be thought of as rings of different sizes (i.e. hours on a clock face). Most importantly remember that division is not necessarily the inverse of multiplication.
- *Floats* - `float`, `double`, `long double`. These are *NOT* the same as  $\mathbb{R}$  (associativity, and even commutativity not guaranteed, multiplicative inverses don't always exist). Programming floats well for numerical problems with large/small numbers is an art form.

## Integer Types - For my 32 bit Windows Box

Type	Min	Max
short	-32768	32767
unsigned short	0	65535
int	-2147483648	2147483647
unsigned int	0	4294967295
long	-2147483648	2147483647
unsigned long	0	4294967295
long long	-9223372036854775808	9223372036854775807
unsigned long long	0	18446744073709551615

For example, here are two bit patterns for `short`:

`1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1` = -1  
`0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1` = 7

(for more information see `<limits.h>`)

## Integer Types(2)

- Two main subtypes *signed* and *unsigned*. Signed types use a sign bit.
- For signed types we, usually, have:
  - minimum value:  $-2^{\text{size}-1}$
  - maximum value:  $2^{\text{size}-1} - 1$
- For unsigned types we have:
  - minimum value: 0
  - maximum value:  $2^{\text{size}} - 1$
- `short` is often used to conserve memory.
- `int` represents the *native* CPU integer type so is used for speed. (If in doubt use `int`).
- `long` and `long long` are used to maintain accuracy.

## Integer Arithmetic

### Base Operators

The four usual operators are defined  $+$ ,  $-$ ,  $*$  and  $/$ .

### Ring arithmetic

Division is not always the reverse of multiplication:

$1/2=0$ ,  $0*2=0$ .

Also, any result of a computation must lie within the ring, any number outside the range of the current data type will "wrap" around. (i.e. 11am + 3 hours gives 2pm).

### Remainder Operator

The remainder operator  $\%$  is unique to integer types, it acts as expected:  $7\%2 = 1$ .

## Floating Point Numbers (IEEE 754 Standard)

On my machine, a `float` (single precision) looks like:



It consists of three parts, the *sign bit*( $b$ ), the *biased exponent*( $e$ ) and the *fraction*( $f$ ). We break down a number  $x$ :

$$x^{\text{float}} = (-1)^b \times 2^{e-127} \times (1 + f \times 2^{-23}), \quad \begin{matrix} 0 < e < 255 \\ 0 \leq f \leq 2^{23} - 1 \end{matrix}$$

We have three special numbers,  $-\text{Inf}$  ( $-\infty$ ),  $\text{Inf}$  ( $\infty$ ) and NaN (Not a Number).

For `double` (double precision) we have:

$$x^{\text{double}} = (-1)^b \times 2^{e-1023} \times (1 + f \times 2^{-52}), \quad \begin{matrix} 0 < e < 2047 \\ 0 \leq f \leq 2^{52} - 1 \end{matrix}$$

## Floating Point

### Base Operators

As with integers, we have  $+$ ,  $-$ ,  $*$  and  $/$ .

### Floating point code

- It looks like integer code but with a decimal point suffix.
- Scientific notation is achieved with `e`:  
`double speedofLight = 2.997e8;` ( $2.997 \times 10^8$ )

### Float Arithmetic

- Division is not always the reverse of multiplication.
- Operators may not be commutative!

$$A + B + C \neq A + C + B \quad (\text{sometimes})$$

## The `pow(x, y)` function (declared in `<math.h>`)

### Exponentiation

There is no exponentiation operator (e.g.  $\wedge$ ,  $**$ ) in C. Instead we have the following:

$$x^y = \text{pow}(x, y)$$

This assumes  $x$  and  $y$  are of type `double`.

### Beware

The `pow` function is often implemented as:

$$\exp(y * \ln(x))$$

For whole integer powers (i.e.  $x^2$ ), one should perform the multiplication explicitly (`x*x`).

## More Mathematical Functions in `<math.h>`

- Maths functions come with the *ANSI Standard C Library*, which contains many maths functions. To use them we need a: `#include <math.h>`
- Here some example functions:  

<code>sin(x)</code>	<code>asin(x)</code>	<code>sinh(x)</code>	<code>exp(x)</code>
<code>cos(x)</code>	<code>acos(x)</code>	<code>cosh(x)</code>	<code>log(x)</code>
<code>tan(x)</code>	<code>atan(x)</code>	<code>tanh(x)</code>	<code>log10(x)</code>
<code>sqrt(x)</code>	<code>atan2(x,y)</code>	<code>pow(x,y)</code>	<code>fabs(x)</code>

  
(all the trigonometric functions use radians!)

## Commenting C Programs

There are two ways of commenting files in C.

### Traditional Way

Anything between `/*` and `*/` is a comment, i.e.

```
/* Hello World! */  
and,  
/* This function is used to compute the  
roots of a quadratic equation */
```

### C++ Style

These are single line only, anything after `//` is a comment, i.e.

```
int c = 3; // set c to 3
```

Technically, C++ style comments aren't in the C standard. (But they are ubiquitous to C code anyway).

## Variable Names

### From K&R

"... Is a sequence of letters and digits. The first character must be a letter; the underscore `_` counts as a letter. Upper and lower case letters are different."

- Punctuation or any other symbols are not allowed in variable names.
- The modern C standard discourages the use of an underscore as the first character of a variable name.

## Simple Logical Expressions

- Are used to carry out branches (`if` statement) and loops (such as `for`, and `while`).
- Evaluate to either *true* (non-zero `int`) or *false* (zero).

### Logical Operators

<code>x</code>	<code>&gt;</code>	<code>y</code>	is <code>x</code> greater than <code>y</code> ?
<code>x</code>	<code>&gt;=</code>	<code>y</code>	is <code>x</code> greater than or equal to <code>y</code> ?
<code>x</code>	<code>&lt;</code>	<code>y</code>	is <code>x</code> less than <code>y</code> ?
<code>x</code>	<code>&lt;=</code>	<code>y</code>	is <code>x</code> less than or equal to <code>y</code> ?
<code>x</code>	<code>==</code>	<code>y</code>	is <code>x</code> equal to <code>y</code> ?
<code>x</code>	<code>!=</code>	<code>y</code>	is <code>x</code> different to <code>y</code> ?



## Compound Logical Expressions

We can create compound logical expressions using the following operators:

- `||` is a *logical or*. `le1 || le2` returns false if both `le1` and `le2` are false and true otherwise.
- `&&` is a *logical and*. `le1 && le2` returns true if and only if both `le1` and `le2` are true.
- `!` is a *logical not*. `!le1` returns the opposite of `le1`.

Here are two identical examples:

- `(x < 100) && (x%2 == 0)`
- `(x < 100) && !(x%2)`

## Flow Control - if

Executes block(s) of code depending on the evaluation of a logical expression.

### Simple if

```
if (logical expression) {statements;}
```

### if, else if, else

```
if (logical expression)
{statements;}
else if (logical expression)
{statements;}
else if (logical expression)
{statements;}
else
{statements;}
```

## Flow Control - while

A `while` loop is used to repeatedly execute code as long as a logical expression is true.

### Structure

```
while (logical expression)
{ statements ;}
```

- If *logical expression* is false, then the *statements* are never executed.

## Flow Control - do {} while ()

We place the *logical expression* after the *statements* giving us:

### Structure

```
do {statements;}
while (logical expression)
```

- The *statements* are executed at least once.

### do while or while?

Generally I prefer `while` over `do while`, as it forces me to initialise variables properly.

## Flow Control - for loop

```
for ( start expression ;
      logical expression ;
      step expression )
{ statements ; }
```

- Print out ten numbers:

```
for (x=0; x < 10; x = x + 1)
    printf("x = %d\n", x);
```

- Keep looping indefinitely (printing out dots)

```
for (;;) printf(".");
```

## Flow Control - switch - case

We can selectively execute code based on a value, using the following:

```
switch (integer_statement) {
    case integer_value1: statements1; break;
    case integer_value2: statements2; break;
    case integer_value3:
    case integer_value4: statements3; break;
    default: statements4; break; }
```

- Execution starts at either one of the `case`'s or at `default`.
- Execution stops at the end `}` or at `break`.
- `case`, `default` and `break` are optional.

## Some Loop Control Features

Execution of code inside a loop (`do`, `while`, `for`) can be manipulated by the following statements.

`break;`

Break out of the current loop. Any statements in the loop following the `break` are ignored and the loop condition automatically evaluates to false, ending the loop.

`continue;`

Jump to the end of the current loop (effectively ignoring everything below the `continue` statement. Whether or not the loop continues executing depends on the loop condition.

## scanf ( ) - Reading Data from Standard Input

For two variables A and B, both of type `double`, we use:

```
scanf("%lf %lf", &A, &B);
```

- where the % represent *format specifiers*

### Format Specifiers

Consist of a %, a numerical width specification and a field code:

d	int	g	float (general form)
u	unsigned int	lf	double (fixed form)
f	float (fixed form)	le	double (exponential form)
e	float (exponential form)	lg	double (general form)

- and the & represents the *address* of the variable in memory. This is known as a *pointer reference operator*.

## Why the &A in scanf ( ) ?

- Functions in C can return only one value.
- Sometimes we want more than one value to change.
- If we tell `scanf` *where* the variables are in memory, `scanf` can change them itself.

The ability to manipulate memory directly is what makes C so powerful. (and potentially dangerous).

## Pointers

A *pointer* is a variable that stores a memory location, they are declared as follows:

```
double * ptrA;
```

### & - Pointer reference operator

Returns the memory address (pointer to) of a variable.

```
double * ptrA = &A;    // ptrA points to A
```

### \* - Pointer de-reference operator

Converts a memory address to a variable:

```
* ptrA = 1.234;        // A is now 1.234
```

## Defining Functions

The C language only provides essential functionality, meaning a lot of functions need to be written yourself. Here are a few general rules for functions:

- Functions cannot define other functions within them.
- An optional single value can be returned.
- All arguments to a functions are passed by value and remain unaffected by the function.
- Passing pointers to functions allows them to “return” multiple variables.

## An Example: Quadratic Equation Solver

As a worked example we write a function to solve the quadratic equation:

$$Ax^2 + Bx + C = 0 \quad A, B, C \in \mathbb{R}$$

Our quadratic solver will:

- Take the three doubles `A`, `B` and `C` as arguments.
- Solve the quadratic and return an `int` signifying to the caller the type of answer available:
  - 1 `A` = 0, we have a linear equation.
  - 0 There are two distinct real roots.
  - 1 We have a pair of complex conjugate roots.
  - 2 Both roots are real and identical.

## The Code

One possible function prototype is:

```
int quad_roots (double A, double B, double C,
               double * r1, double * r2);
```

- The variables A, B and C are unchanged by `quad_roots`.
- We need to return two doubles (the roots of the equation), thus we take in pointers `double *r1` and `double *r2`.
- C90 does not allow for complex number types (C99 does support them), so we have to think a little bit about the complex number case.

## Code Snippet for Calling `quad_roots`

```
...
int main()
{
    double A, B, C, root1, root2;
    int quad_case;
    ...
    quad_case = quad_roots(A, B, C, &root1,
                          &root2);

    switch(quad_case)
    {
        case -1: linear equation
```

## Code Snippet for `quad_roots`

```
int quad_roots(double A, double B, double C,
               double * r1, double * r2)
{
    double d;

    /* linear case */
    if (A == 0.0)
    {
        *r1 = -C/B;
        return -1;
    }

    /* compute the discriminant */
    d = B*B-4.0*A*C;
```

## Declarations vs Definitions

### Function Declarations

These tell the compiler about the *existence* of a function, which then allows us to call it. A declaration ends with a `;`.

```
int quad_roots (double A, double B, double C,
               double * r1, double * r2);
```

### Function Definitions

The code making up the function is supplied to the compiler. A function can only be defined once. A definition contains braces `{` and `}`:

```
int quad_roots (double A, double B, double C,
               double * r1, double * r2)
{ ... }
```