

## Computing in C for Science

### Lecture 3 of 5

Dr. Steven Capper

[steven.capper99@imperial.ac.uk](mailto:steven.capper99@imperial.ac.uk)

<http://www2.imperial.ac.uk/~sdc99/ccourse/>

30<sup>th</sup> November 2011

Imperial College  
London

Steven Capper

C for Science - Lecture 3

## Common Problems

### ❶ Misuse of Power!

$x^2$  should always be written `x*x`, *NOT* `pow(x, 2.0)`.

$\sqrt{x}$  should be written `sqrt(x)`, *NOT* `pow(x, 0.5)`.

### ❷ Integer Division

Remember a fraction such as  $1/3$  is equal to zero. To get a floating point fraction, this should be rewritten `1.0/3.0`.

Steven Capper

C for Science - Lecture 3

## Input and Output in C - Streams

- All I/O in C is accomplished via file *streams*.
- This stems from C's UNIX roots (every device is a file).
- We have already seen `printf`, formatted output to console.
- C90 defines three streams which are initialised by default when a program starts.

|                     |                                      |
|---------------------|--------------------------------------|
| <code>stdin</code>  | Input from the keyboard. (read only) |
| <code>stdout</code> | Text console. (write only)           |
| <code>stderr</code> | Another text console. (write only)   |

Steven Capper

C for Science - Lecture 3

## Why have `stderr`?

### Stream Redirection

- UNIX (and now most other operating systems) allow for a C program's streams to be *redirected* at the command line.
- This is completely transparent to the C program (it just reads and writes data oblivious to the source).
- One can redirect `stdout` using `(>)` as follows:  

```
myprogram > output.txt
```
- `stderr` can be redirected using `(2>)`:  

```
myprogram 2> errorlog.txt
```
- `stdin` can be redirected using `(<)`:  

```
myprogram < input.txt
```

Having a separate *error stream*, allows for important messages to be filtered from data output.

Steven Capper

C for Science - Lecture 3

## printf and scanf - formatted output and input

These are automatically connected to `stdout` and `stderr` respectively. The more generalised functions are `fprintf` and `fscanf`. As an example:

```
fprintf(stdout, "Text to stdout...\n");
fprintf(stderr, "Text to stderr...\n");
```

We can also read text from a stream:

```
fscanf(stdin, "%lf", &ptrDouble);
```

- Additional streams can be opened using `fopen`.

## fopen - open a file stream

```
FILE * stream = fopen(char * filename, char * mode);
```

Working backwards, `mode` is a string telling C how to open the file:

| mode | meaning   |
|------|---|
| "r"  | Open filename for reading. The file must exist or NULL is returned.   |
| "w"  | Open filename for writing, starting from the beginning of the file. The file will be created if it doesn't already exist. Any old data will be overwritten. |
| "a"  | Open filename for writing, starting from the end of the file. File will be <i>appended</i> if it doesn't exist and created otherwise.                       |
| "r+" | Open filename for reading and writing, starting from the beginning. If the file doesn't exist, NULL is returned.  |
| "w+" | Open filename for reading and writing, starting from the beginning. If the file doesn't exist it's created.   |
| "a+" | Open filename for reading and writing (append if exists).   |

## fopen - II

```
FILE * stream = fopen(char * filename, char * mode);
```

### filename

Must be a legal operating system file name. A safe option would be something similar to:

```
"results.txt"
```

### Absolute Filenames

We can specify a full directory path to a filename:

```
data = fopen("C:\\TEMP\\mydata.dat", "w"); /* Windows */
data = fopen("/tmp/mydata.dat", "w");      /* UNIX */
```

### Relative Filenames

It is much safer to omit the full path:

```
data = fopen("mydata.dat", "w"); /* works on most os' */
```

## fopen - III

```
FILE * stream = fopen(char * filename, char * mode);
```

- Here `stream` is a pointer to a `FILE` data type.
- `stdin`, `stdout` and `stderr` are also pointers (which are opened automatically for us).
- Sometimes it is not possible to open a file (it may not exist, or belong to another user etc...), in this case a special pointer value `NULL` (which is equal to zero) is returned.

### NULL pointers

A `NULL` pointer is a special value, usually signalling failure. This can be checked for:

```
stream = fopen("illegal/\\filename.txt", "w");
if (!stream)
{
    fprintf(stderr, "Unable to open file\n");
    ...
}
```

## fclose - Closing a file stream

- Any file streams `fopen`'ed should be closed using:

```
fclose(stream);
```

- File read/writes may go to an internal buffer before they go to the disk (to speed things up).
- An explicit file close ensures that any pending data is written to disk.
- For the lazy there is:

```
fcloseall();
```

- Do *NOT* call `fclose` on any of the following:  
`stdin, stdout & stderr`

## Moving around a file - `rewind`, `ftell` and `fseek`

```
rewind(stream);
```

Start the next read/write from the beginning of the file.

```
fpos = ftell(stream);
```

Get an `int` telling us how far we are in the file (from the beginning).

```
fseek(stream, fpos, SEEK_SET);
```

Move `fpos` bytes into the file (counted from the beginning).

```
fseek(stream, fpos, SEEK_CUR);
```

Move `fpos` bytes forward into file (from current position).

```
fseek(stream, fpos, SEEK_END);
```

Move `fpos` bytes from the end of the file.

## File Types

Files in C can contain data in two formats:

- As text, where each byte is interpreted as an ASCII character mapping.
- Or as raw *binary* data, the meaning of which is left to the program to determine.

These files are referred to as text files and binary files respectively. All the I/O covered so far applies to text files.

Each format has its advantages and disadvantages.

## Text versus Binary Files

### Text Files

- Pro** - Can be read by almost any program on almost any computer. (And by humans!).
- Con** - Inefficient - Numbers can take up as much as twice the space as necessary.

### Binary Files

- Pro** - Accurate and efficient.
- Con** - Incompatible - Binary data can be very difficult to read by other programs, and incredibly difficult between processor architectures.

## Handling Binary Files

Binary files are opened using `fopen`, we need to add "b" to the file mode though:

```
stream = fopen("data.dat", "rb");
```

Data is read from and written to a binary file using `fread` and `fwrite`. For example:

```
double values[10] = {0.0};
/* we have "stream" an open binary file */
read = fread(values, sizeof(double), 10, stream);
if (read != 10)
{ read failed }
```

Please see `<stdio.h>` for more details.

## Permanent versus Temporary Files

- Sometimes we may not wish for a file to be permanent.
- Temporary files provide a means for storing intermediate computations.
- We open a temporary file using the `tmpfile()` function.

### Example

```
stream = tmpfile();
if (!stream)
{
    /* handle failure */
}
/* fread, fwrite, fprintf, fscanf */
fclose(stream);
```

## Pointers and Arrays

We saw the connection between arrays and pointers in the last lecture. Given the array:

```
double ar [4] = {1.0, 2.0, 3.0, 4.0};
```

The elements of this array can be accessed as follows:

```
*(ar) == ar[0] == 1.0
*(ar+1) == ar[1] == 2.0
*(ar+2) == ar[2] == 3.0
```

Pointer referencing is also supported:

```
ar == &ar[0]
(ar+1) == &ar[1]
(ar+2) == &ar[2]
```

## Allowed Pointer Operations

|                 |                                 |
|-----------------|---------------------------------|
| Declaration:    | <code>double * pA, * pB;</code> |
| Assignment:     | <code>pA = &amp;var;</code>     |
| Increment:      | <code>pA = pA + 1;</code>       |
| Decrement:      | <code>pA = pA - 1;</code>       |
| Difference:     | <code>gap = pA - pB;</code>     |
| Comparison:     | <code>if(pA == pB)</code>       |
| De-referencing: | <code>*pA = val;</code>         |

## Fixed Size Two-Dimensional Arrays

We can declare arrays of dimension higher than one, as follows:

```
double a[2][3] = {{1.0, 2.0, 3.0},
                 {2.0, 3.0, 4.0}};
```

Where the elements of `a` are denoted as expected:

|                      |                      |                      |
|----------------------|----------------------|----------------------|
| <code>a[0][0]</code> | <code>a[0][1]</code> | <code>a[0][2]</code> |
| <code>a[1][0]</code> | <code>a[1][1]</code> | <code>a[1][2]</code> |

To access the top left element:

```
myVal = a[0][0]; /* equal to 1.0 */
```

## Fixed Size Two-Dimensional Arrays II

```
#include <stdio.h>
#define COLS 3

void printArray(int matrix[][COLS], int rows)
{
    int i, j;
    for (i = 0; i < rows; i++)
    {
        for (j = 0; j < COLS; j++)
            printf("%d ", matrix[i][j]);
        printf("\n");
    }
}

int main()
{
    int matrix[2][COLS] = {{1, 2, 3},{4, 5, 6}};
    printArray(matrix, 2);
    return 0;
}
```

## Fixed Size Two-Dimensional Arrays III

We have seen an example of a two dimensional array:

|                      |                      |                      |
|----------------------|----------------------|----------------------|
| <code>a[0][0]</code> | <code>a[0][1]</code> | <code>a[0][2]</code> |
| <code>a[1][0]</code> | <code>a[1][1]</code> | <code>a[1][2]</code> |

- In memory it is arranged as follows:  

|                      |                      |                      |                      |                      |                      |
|----------------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| <code>a[0][0]</code> | <code>a[0][1]</code> | <code>a[0][2]</code> | <code>a[1][0]</code> | <code>a[1][1]</code> | <code>a[1][2]</code> |
|----------------------|----------------------|----------------------|----------------------|----------------------|----------------------|

  
(i.e. as a one dimensional array).
- Fixed size arrays are very inflexible as they require the dimensions to be "hard coded".
- They are allocated from the stack thus large arrays may cause problems.
- *Dynamically allocated* arrays overcome these restrictions.

## Dynamic Allocation

### The heap

- C has set aside a special block of memory known as the *heap*.
- Memory can be requested from the heap.
- If memory is available, it is provided, otherwise a request is made to the operating system.
- The implementation details of the heap are "undefined".

### sizeof

- We need to know how much memory to request.
- The `sizeof` keyword returns the size (in bytes), of a variable or data type.

## How to Borrow Memory

We use the `malloc` function (defined in `<stdlib.h>`) to request memory and `free` to return it to the heap.

```
#include <stdio.h>
#include <stdlib.h>
#define NCOUNT 10
int main()
{
    int * numbers, loop;
    numbers = (int *) malloc(sizeof(int)*NCOUNT);
    if (!numbers)
    {
        fprintf(stderr, "Unable to allocate memory\n");
        return -1;
    }
    for (loop = 0; loop < NCOUNT; loop++)
        numbers[loop] = loop;

    for (loop = 0; loop < NCOUNT; loop++)
        printf("numbers[%d] = %d\n", loop, numbers[loop]);

    free(numbers);
    return 0;
}
```

Steven Capper

C for Science - Lecture 3

## How to Borrow Memory - Details

```
void * malloc(size_t Size);
void free(void * Memory);
```

### What is `void *`?

- `void *` is a pointer to an unknown data type, it consists only of an address.
- We can explicitly cast from `void *` to any other pointer type.
- Any pointer type can be implicitly casted to a `void *`.
- A value of `NULL` usually means that a function encountered an error.

### What is `size_t`?

`size_t` is defined to be an unsigned integer type large enough to hold numbers returned by `sizeof`.

Steven Capper

C for Science - Lecture 3

## How to Borrow Memory - More Details

Two more popular memory management functions are:

|                      |   |
|----------------------|---|
| <code>realloc</code> | Modifies the size of a memory block.        |
| <code>calloc</code>  | Allocates memory, and sets every byte to 0. |

### Some common heap problems

- Allocating memory is slow and should be done as few times as possible.
- Memory allocation can fail, every `malloc/calloc` should be checked (or at the very least asserted).
- Every `malloc/calloc` should have a corresponding `free`, memory is *leaked* if it's not freed after use.

Steven Capper

C for Science - Lecture 3

## 1-D Example: Fibonacci (again!)

```
#include <stdlib.h>
#include <stdio.h>

int main()
{
    double * fibs;
    unsigned int nfibs=0, loop;
    while (nfibs < 2)
    {
        printf("How many Fibonacci numbers are needed (>1)?\n");
        scanf("%u", &nfibs);
    }

    fibs = (double *) malloc(sizeof(double)*nfibs);
    if (!fibs) /* malloc failed? */
    {
        fprintf(stderr, "Unable to allocate memory!\n");
        return -1;
    }

    fibs[0] = 1.0; fibs[1] = 1.0;
    for (loop = 2; loop < nfibs; loop++)
        fibs[loop] = fibs[loop-1]+fibs[loop-2];

    for (loop = 0; loop < nfibs; loop++)
        printf("fib[%u] = %lg\n", loop, fibs[loop]);

    free(fibs);
    return 0;
}
```

Steven Capper

C for Science - Lecture 3

## Pointers to Pointers: \*\*ptr

- Pointers in C are very flexible, to the extent that we can form pointers to pointers!
- This is known more formally as the “level of indirection”.
- Tensors of arbitrary rank (Fortran 95 is limited to 7) can be formed easily in C.
- The most useful example in C being matrices.

### Matrices in C

Matrices in C are commonly represented by the `double **` type. This means “a pointer to a pointer to a `double`”.

## Constructing Matrices with Pointers

```
double ** makeMatrix(unsigned int rows, unsigned int cols)
{
    unsigned int i;
    double ** matrix;

    matrix = (double **) malloc(rows * sizeof(double *));
    if (!matrix) return NULL; /* failed */

    for (i = 0; i < rows; i++)
    {
        matrix[i] = (double *) malloc(cols*sizeof(double));
        if (!matrix[i])
            return NULL; /* lazy, we should really free
                           all the memory allocated above */
    }

    return matrix;
}
```

## Accessing Matrix Elements

### Usage pattern for `makeMatrix`

```
double ** matrix = makeMatrix(rows, cols);
for (i=0; i < rows; i++)
    for (j=0; j < cols; j++)
        matrix[i][j] = 0.0;
free the matrix
```

- Accessing the dynamically allocated array looks identical to the fixed size ones, but “under the hood” things are a little different:

```
matrix[row][col] = (*(matrix + row) + col)
```

- The `makeMatrix` code on the previous slide contained a lot of `malloc` statements, is there a better way to allocate a matrix? (yes!)

## A Better Way of Allocating Matrices

```
double ** allocMatrix(unsigned int rows, unsigned int cols)
{
    double ** matrix;
    unsigned int i;

    matrix = (double **) malloc (rows*sizeof(double *));
    if (!matrix) return NULL; /* failed */

    matrix[0] = (double *) malloc (rows*cols*sizeof(double));
    if (!matrix[0])
    {
        free(matrix); /* we don't need matrix any more */
        return NULL; /* failed */
    }

    for (i = 1; i < rows; i++)
        matrix[i] = matrix[i-1] + cols;

    return matrix;
}
```

## Why is allocMatrix Better?

- allocMatrix only uses 2 mallocs whilst, makeMatrix uses cols + 1.
- Meaning there are fewer points of failure (we only check two pointers for NULL),
- It is much easier to free a matrix allocated with the allocMatrix function, all we need to do is:

```
void freeMatrix(double ** matrix)
{
    free(matrix[0]);
    free(matrix);
}
```

## Case Study: Matrix Addition

Let's define some utility functions to:

- 1 Allocate memory for the matrix (allocMatrix) - **done**,
- 2 Free a matrix (freeMatrix) - **done**,
- 3 Print a matrix (printMatrix),
- 4 Create a random matrix (randomMatrix),
- 5 Add matrices together (addMatrices)

We drive all these functions using a main function.

## printMatrix and randomMatrix

```
void printMatrix(double ** matrix, unsigned int rows,
                unsigned int cols)
{
    unsigned int i, j;

    for (i = 0; i < rows; i++)
    {
        for (j = 0; j < cols; j++)
            printf("%8.5lf ", matrix[i][j]);
        printf("\n");
    }
}

void randomMatrix(double ** matrix, unsigned int rows,
                unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++)
        for (j = 0; j < cols; j++)
            matrix[i][j] = (double)rand()/RAND_MAX;
}
```

## addMatrices

```
void addMatrices(double ** matrixA, double ** matrixB,
                double ** matrixR,
                unsigned int rows, unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++)
        for (j = 0; j < cols; j++)
            matrixR[i][j] = matrixA[i][j] + matrixB[i][j];
}
```



## The main function

```
int main()
{
    unsigned int rows, cols;
    double ** matrixA, ** matrixB, **matrixC;
    printf("Enter rows cols: ");
    scanf("%u %u", &rows, &cols);

    matrixA = allocMatrix(rows, cols);
    matrixB = allocMatrix(rows, cols);
    matrixC = allocMatrix(rows, cols);

    if (!matrixA || !matrixB || !matrixC)
    {
        /* a little lazy, but it does the job */
        fprintf(stderr, "Unable to allocate matrices!\n");
        return -1;
    }

    randomMatrix(matrixA, rows, cols); randomMatrix(matrixB, rows, cols);
    addMatrices(matrixA, matrixB, matrixC, rows, cols);

    printf("\n\nmatrix A = \n");
    printMatrix(matrixA, rows, cols);
    printf("\n\nmatrixB = \n");
    printMatrix(matrixB, rows, cols);
    printf("\n\nmatrixA + matrixB = \n");
    printMatrix(matrixC, rows, cols);

    freeMatrix(matrixC); freeMatrix(matrixB); freeMatrix(matrixA);
    return 0;
}
```

## Results

Enter rows cols: 4 4

matrix A =

|         |         |         |         |
|---------|---------|---------|---------|
| 0.00125 | 0.56359 | 0.19330 | 0.80874 |
| 0.58501 | 0.47987 | 0.35029 | 0.89596 |
| 0.82284 | 0.74660 | 0.17411 | 0.85894 |
| 0.71050 | 0.51353 | 0.30399 | 0.01498 |

matrixB =

|         |         |         |         |
|---------|---------|---------|---------|
| 0.09140 | 0.36445 | 0.14731 | 0.16590 |
| 0.98853 | 0.44569 | 0.11908 | 0.00467 |
| 0.00891 | 0.37788 | 0.53166 | 0.57118 |
| 0.60176 | 0.60717 | 0.16623 | 0.66305 |

matrixA + matrixB =

|         |         |         |         |
|---------|---------|---------|---------|
| 0.09265 | 0.92804 | 0.34062 | 0.97464 |
| 1.57353 | 0.92557 | 0.46937 | 0.90063 |
| 0.83175 | 1.12449 | 0.70577 | 1.43013 |
| 1.31227 | 1.12070 | 0.47023 | 0.67803 |