

Computing in C for Science

Lecture 4 of 5

Dr. Steven Capper

steven.capper99@imperial.ac.uk

<http://www2.imperial.ac.uk/~sdc99/ccourse/>

7th December 2011

Imperial College
London

Steven Capper

C for Science - Lecture 4

Cryptic C

There are shortcuts in the C language that allow for concise code.

- 1 Incrementing by 1. Pre-increment, and post-increment.
++i Increment i by 1, then use it.
i++ Use i, then increment it by 1.
- 2 Increment by another variable.
Normal code: `sum = sum + v[i];`
Terse code: `sum += v[i];`

An example:

```
for (i=0; i < N; i++)  
    sum += v[i];
```

Steven Capper

C for Science - Lecture 4

More Shorthand

```
--i;          decrement i by 1.  
sum -= v[i];  means sum = sum - v[i];  
sum *= v[i];  means sum = sum * v[i]; Other  
sum /= v[i];  means sum = sum / v[i];  
sum %= 2;     means sum = sum % 2;  
operators can also be abbreviated this way.
```

Inline if

The following code:

```
if (r1 > r2) maxr = r1; else maxr = r2;
```

can be abbreviated:

```
maxr = (r1 > r2) ? r1 : r2;
```

Steven Capper

C for Science - Lecture 4

Matrices - Details

For matrices allocated in the previous lecture

```
matrix        points to a matrix  
matrix[0]     points to the first row of the matrix  
matrix[0][0]  The top left element of the matrix
```

Which have the following types:

```
matrix        double **  
matrix[0]     double *  
matrix[0][0]  double
```

Steven Capper

C for Science - Lecture 4

Higher Dimensional Arrays

```
double *** alloc3Tensor(unsigned int dim)
{
    unsigned int i, j;
    double *** tensor;
    tensor = (double ***) malloc(dim*sizeof(double **));
    if (!tensor) return NULL;
    tensor[0] = (double **) malloc(dim*dim*sizeof(double *));
    if (!tensor[0])
    {
        free(tensor);
        return NULL;
    }
    tensor[0][0] = (double *) malloc(dim*dim*dim*sizeof(double));
    if (!tensor[0][0])
    {
        free(tensor[0]); free(tensor);
        return NULL;
    }
    for (i = 1; i < dim; i++)
        tensor[0][i] = tensor[0][i-1] + dim;
    for (i = 1; i < dim; i++)
    {
        tensor[i] = tensor[i-1] + dim;
        tensor[i][0] = tensor[i-1][0] + dim * dim;
        for (j = 1; j < dim; j++)
            tensor[i][j] = tensor[i][j-1] + dim;
    }
    return tensor;
}
```

More Multi-Index

- The 3D array behaves as expected:

```
tensor[i][j][k] = 0.5;
```

- The array can be freed with:

```
void free3Tensor(double *** tensor)
{
    free(tensor[0][0]);
    free(tensor[0]);
    free(tensor);
}
```

More Elaborate Data Structures

- In C we are able to manipulate data directly, this has allowed us to partition a contiguous array of memory into matrix rows.
- We needn't restrict ourselves to structures where the number of elements per row is constant, one example is Pascal's triangle:

```
    1
   1 2 1
  1 3 3 1
 1 4 6 4 1
   ⋮
```

```
int main()
{
    unsigned int size = 10, r, c;
    int ** pasc;
    pasc = (int **) malloc(size*sizeof(int));
    pasc[0] = (int *) malloc(size*(size+1)*sizeof(int)/2);
    /* check the mallocs */
    for (r=1; r < size; r++) pasc[r] = pasc[r-1]+r;
    pasc[0][0] = 1;
    for (r = 1; r < size; r++)
    {
        for (c = 1; c < r; c++)
            pasc[r][c] = pasc[r-1][c-1]+pasc[r-1][c];
        pasc[r][0] = 1; pasc[r][r] = 1;
    }
    for (r=0; r < size; r++)
    {
        for (c = 0; c < r+1; c++)
            printf("%3d ", pasc[r][c]);
        printf("\n");
    }
    free(pasc[0]);
    free(pasc);
    return 0;
}
```

Pascal's Triangle

- The program outputs:

```
1
1 1
1 2 1
1 3 3 1
1 4 6 4 1
1 5 10 10 5 1
1 6 15 20 15 6 1
1 7 21 35 35 21 7 1
1 8 28 56 70 56 28 8 1
1 9 36 84 126 126 84 36 9 1
```

- Whilst in memory, the data structure looks like:

1	1	1	1	2	1	1	3	3	1	...
---	---	---	---	---	---	---	---	---	---	-----

Custom Data Types

Recall, when dealing with files we used:

```
FILE * file;
file = fopen ("myfile.txt", "r");
```

FILE is in fact a custom data type, with its own size:

```
#include <stdio.h>
```

```
int main()
{
    printf("sizeof(FILE) = %d\n", sizeof(FILE));
    return 0;
}
```

Excerpt from MS - <stdio.h>

(From Visual Studio 2008)

```
struct _iobuf {
    char *_ptr;
    int _cnt;
    char *_base;
    int _flag;
    int _file;
    int _charbuf;
    int _bufsiz;
    char *_tmpfname;
};
typedef struct _iobuf FILE;
```

typedef Structures : Custom Data Types

- FILE is a custom data type defined in <stdio.h>.
- It is comprised of elements of different, known, types.
- Each element of the FILE structure has its own name.

Let's consider a structure of our own, one of the simplest examples is a complex number:

```
typedef struct
{
    double real;
    double imag;
} complex;
```

C99

C99 fully supports its own complex type.

typedef Structures - II

Definitions of structures take the following form:

```
typedef struct
{
    elementType elementName;
    elementType elementName;
    :
} structureTypeName ;
```

Handling Structures

- 1 A structure may be an argument to a function.
- 2 A function may return a structure.
- 3 A pointer may point to a structure.
- 4 Structures are referenced as normal: `&name`.
- 5 Elements of a structure are references as: `&name.element`.
- 6 A pointer to a structure may be passed to a function.
- 7 If `p` is a pointer to a structure, then `p->member` allows us to access it's members.

Example Function Applying to Structures

```
#include <stdio.h>

typedef struct
{
    double real;
    double imag;
} complex;

void printComplex(complex * mc)
{
    printf("%lg + %lgi\n", mc->real, mc->imag);
}

int main()
{
    complex c1 = {1.0, 0.5}; /* assignment at declaration */
    printComplex(&c1);       /* pass pointer to struct */
    c1.real = 10.0;          /* piecewise assignment */
    printComplex(&c1);
    return 0;
}
```

More Structures

- Passing structures via pointer is usually more efficient.
- Assignment can be done at declaration or after.
- In C we are not allowed to overload operators, so the following won't work:
$$c1 = c2 + c3; \text{ (where } c1 \text{ and } c2 \text{ are complex)}$$
so we need to do something like:
$$\text{complexAdd}(\&c1, \&c2, \&c3);$$
- Arrays of `structs` are allowed (so matrices can consist of complex numbers for example).
- Structures can contain structures as elements.

Complex Number Support in C/C++

C++

Complex numbers are supported in C++ as `classes`, also the operators are overloaded properly too meaning any code which uses them will be concise. (look in `<complex>`).

C99

Complex numbers are supported as a native data type (not `struct`) in C99. Unfortunately not many compilers support this. GNU C, fully supports complex numbers (see `<complex.h>`).

C90

Complex number support in C90 is non-existent. I would recommend either third party libraries or a switch to C99/C++ for heavy complex number use.

Another Example `struct` - from `<time.h>`

```
struct tm
{
    int tm_sec;      /* seconds after the minute - [0,59] */
    int tm_min;      /* minutes after the hour - [0,59] */
    int tm_hour;      /* hours since midnight - [0,23] */
    int tm_mday;      /* day of the month - [1,31] */
    int tm_mon;       /* months since January - [0,11] */
    int tm_year;       /* years since 1900 */
    int tm_wday;       /* days since Sunday - [0,6] */
    int tm_yday;       /* days since January 1 - [0,365] */
    int tm_isdst;      /* daylight savings time flag */
};
```

Compilation

- As seen in the first lecture, C programs are *compiled* into a low level (machine specific language).
- This language is called *assembly language*.
- On PCs and Macs Intel x86 assembler is ubiquitous.
- Most C compilers allow you to view the assembler that they generate.

addMatrices revisited - C code

From the last lecture we saw a function to add two matrices together:

```
void addMatrices(double ** matrixA, double ** matrixB,
                 double ** matrixR,
                 unsigned int rows, unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++)
        for (j = 0; j < cols; j++)
            matrixR[i][j] = matrixA[i][j] + matrixB[i][j];
}
```

When we compile this, we get...

addMatrices revisited - when compiled...

```

_addMatrices PROC
; 35 : {
        push    ebp
        mov     ebp, esp
        sub     esp, 216
        push    ebx
        push    esi
        push    edi
        lea     edi, DWORD PTR [ebp-216]
        mov     ecx, 54
        mov     eax, -858993460
        rep stosd
; 36 :     unsigned int i, j;
; 37 :     for (i = 0; i < rows; i++)
        mov     DWORD PTR _i$[ebp], 0
        jmp     SHORT $LN6@addMatrice
$LN5@addMatrice:
        mov     eax, DWORD PTR _i$[ebp]
        add     eax, 1
        mov     DWORD PTR _i$[ebp], eax
$LN6@addMatrice:
        mov     eax, DWORD PTR _i$[ebp]
        cmp     eax, DWORD PTR _rows$[ebp]
        jae     SHORT $LN4@addMatrice
; for (j = 0; j < rows; j++)
        mov     DWORD PTR _j$[ebp], 0
        jmp     SHORT $LN3@addMatrice
$LN2@addMatrice:
        mov     eax, DWORD PTR _j$[ebp]
        add     eax, 1
        mov     DWORD PTR _j$[ebp], eax
        $LN3@addMatrice:
        mov     eax, DWORD PTR _j$[ebp]
        cmp     eax, DWORD PTR _rows$[ebp]
        jae     SHORT $LN1@addMatrice
        $LN4@addMatrice:
        pop     edi
        pop     esi
        pop     ebx
        mov     esp, ebp
        pop     ebp
        ret     0
_addMatrices ENDP

```

Steven Capper

C for Science - Lecture 4

Optimisation

- A few lines of C becomes > 30 lines of assembler (only three of which are actually floating point instructions!).
- It is possible to write a much smaller assembler routine by hand ≈10-20 instructions long.
- This would run ≈3 times quicker than the C compiled routine (this is a general rule of thumb).
- Any custom assembly code would only target a very specific chip, however.

Steven Capper

C for Science - Lecture 4

Optimisation II

- Rather than rewrite the assembly code, it is easier to ask the C compiler to perform code optimisation itself.
- By default C will compile the code as it appears (the exact order of operations is preserved etc), this aids debugging.
- C compilers can be told to optimise their code in the following ways:

MSVC Project configuration options can be set, defaults in “Release” build do a good job.

gcc The -O command line flags influence optimisation, -O0 means “off” whilst -O3 means “extremely aggressive”.

Steven Capper

C for Science - Lecture 4

Optimisation Example - Visual Studio 2008

Debug Build

```

$LN2@addMatrice:
mov     eax, DWORD PTR _j$[ebp]
add     eax, 1
mov     DWORD PTR _j$[ebp], eax
$LN3@addMatrice:
mov     eax, DWORD PTR _j$[ebp]
cmp     eax, DWORD PTR _rows$[ebp]
jae     SHORT $LN1@addMatrice
;matrixR[i][j] = matrixA[i][j]+matrixB[i][j];
mov     eax, DWORD PTR _i$[ebp]
mov     ecx, DWORD PTR _matrixA$[ebp]
mov     edx, DWORD PTR [ecx+eax*4]
mov     eax, DWORD PTR _i$[ebp]
mov     ecx, DWORD PTR _matrixB$[ebp]
mov     eax, DWORD PTR [ecx+eax*4]
mov     ecx, DWORD PTR _j$[ebp]
mov     esi, DWORD PTR _j$[ebp]
fld     QWORD PTR [edx+ecx*8]
fadd    QWORD PTR [eax+esi*8]
mov     edx, DWORD PTR _i$[ebp]
mov     eax, DWORD PTR _matrixR$[ebp]
mov     ecx, DWORD PTR [eax+edx*4]
mov     edx, DWORD PTR _j$[ebp]
fstp    QWORD PTR [ecx+edx*8]
jmp     SHORT $LN2@addMatrice

```

Release Build

```

;matrixR[i][j] = matrixA[i][j]+matrixB[i][j];
fld     QWORD PTR [edx+eax]
fadd    QWORD PTR [eax]
fstp    QWORD PTR [esi+eax]
add     eax, 8
dec     ebx
jne     SHORT $LN3@addMatrice

```

Steven Capper

C for Science - Lecture 4

Bits and Bytes

Bytes

Smallest *addressable* unit of memory, each byte is composed of eight bits.

Bits

These are the smallest units of computer memory, each bit can be either 0 or 1.

- Addressing bits individually requires some extra operations to be carried out.
- There are good reasons for accessing data at the bit-level however.

Efficient Data Packing

Given a 32 million base pair chromosome

It will require: ≈ 64 megabytes to store as `short`
 ≈ 32 megabytes to store as `char` (next lecture)
 ≈ 8 megabytes to store as bit data.
(i.e. a 2-4 year research lead if following Moore's Law).

Computer Graphics

Given a monochrome print image of 2400 dpi rendered over 80 square inches gives $\approx 500,000,000$ dots. This takes up:

≈ 1 gigabyte if using `short`
 ≈ 64 megabytes if using bits.

Bit Manipulation Friendly Data Types

As seen before the unsigned integer data types have values ranging from 0 to $2^n - 1$ where n is the number of bits in the data type. Some examples (for my machine):

Data Type	n
unsigned short	16
unsigned int	32
unsigned long int	64
unsigned long long int	128

Unsigned data types are also desirable for accessing array indices as they can never be negative.

How to Get Them In and Out of The Computer

- They can be read using `scanf` and `"%u"`, `"%lu"` or `"%Lu"`.
- They can be printed using `printf` and `"%u"` or `"%Lu"`.
- We can output to octal (base 8 numbers) using `printf` and `"%o"`.
- Also we can output to hexadecimal (base 16, 0-9 and a-f) using `printf` and `"%x"` or `"%X"`.

Example - byte (`char`)

Binary	10101011
Hexadecimal	AB
Decimal	171
Octal	253

Manipulating Bits within an Unsigned Integer

- C can shift all the bits comprising a number a fixed number of places to the left or right.
- Zeros are propagated in to the vacated spaces.
- Bits that shift outside, disappear. (i.e. the shift is not cyclic).
- Bit shifting is accomplished with the `>>` (right) and `<<` (left) operators.

For example:

```
1 << 2 = 4
8 >> 3 = 1
```

Bit shifts are much cheaper than multiplying or dividing by powers of two.

4 Bitwise operators `&`, `|`, `^` and `~`

Assuming 0 is false and 1 is true, we have the following bitwise logical operators.

And Operator(`&`)

N1	0	0	1	1
N2	0	1	0	1
N1 & N2	0	0	0	1

Or Operator(`|`)

N1	0	0	1	1
N2	0	1	0	1
N1 N2	0	1	1	1

Exclusive Or(`^`)

N1	0	0	1	1
N2	0	1	0	1
N1 ^ N2	0	1	1	0

Not Operator(`~`)

N1	0	1
~N1	1	0

Case Study: The Sieve of Eratosthenes

3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51
3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51
3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51
3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51
3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51
3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51
3	5	7	9	11	13	15	17	19	21	23	25	27	29	31	33	35	37	39	41	43	45	47	49	51

This gives the primes (we add 2 to the beginning of the list):

2 3 5 7 11 13 17 19 23 29 31 37 41 43 47

The Sieve of Eratosthenes: In C

- When implementing this in C it makes sense to use a bit to indicate 'primeness' of a number.
- The smallest addressable unit of memory in C is the `char` which consists of 8 bits.
- We therefore need to perform *masking* to isolate individual bits.

Masking

We access the j^{th} bit of a variable x as follows:

```
if (x & (1 << j)) Check to see if it's set
x |= (1 << j)      To set the bit
x &= ~(1 << j)     To clear the bit
```


The Sieve of Eratosthenes: Implementation

```
void findPrimes(NUM_TYPE * Prime_List, int max_num)
{
    int current_num = 3;
    NUM_TYPE Mask=1;

    while(current_num*current_num <= max_num)
    {
        int Pnum = current_num/(2 * BITS_PER_NUM);
        int Pbit = (current_num-Pnum * 2 * BITS_PER_NUM)/2;
        /* Current Number is prime so strike out multiples */
        if(~Prime_List[Pnum] & (Mask <<Pbit))
        {
            int strike = current_num * current_num;
            while(strike <= max_num)
            {
                int Snum = strike/(2*BITS_PER_NUM);
                int Sbit = (strike - Snum * 2 * BITS_PER_NUM)/2;
                Prime_List[Snum] |= (Mask << Sbit);
                strike += 2 * current_num;
            }
        }
        current_num += 2;
    }
}
```

The Sieve of Eratosthenes : Printing out the Primes

```
void printPrimes(NUM_TYPE * Prime_List, int max_num)
{
    int i;
    NUM_TYPE Mask=1;

    for(i = 3;i <= max_num;i = i+2)
    {
        int Pnum=i/(2*BITS_PER_NUM);
        int Pbit=(i-Pnum*2*BITS_PER_NUM)/2;
        if (~Prime_List[Pnum] & (Mask <<Pbit))
            printf("%d\n", i);
    }
}
```