

Computing in C for Science

Lecture 4 of 5

Dr. Steven Capper

`steven.capper99@imperial.ac.uk`

<http://www2.imperial.ac.uk/~sdc99/ccourse/>

7th December 2011

Imperial College
London

There are shortcuts in the C language that allow for concise code.

- 1 Incrementing by 1. Pre-increment, and post-increment.

`++i` Increment `i` by 1, then use it.

`i++` Use `i`, then increment it by 1.

- 2 Increment by another variable.

Normal code: `sum = sum + v[i];`

Terse code: `sum += v[i];`

An example:

```
for (i=0; i < N; i++)  
    sum += v[i];
```

More Shorthand

`--i;` decrement `i` by 1.
`sum -= v[i];` means `sum = sum - v[i];`
`sum *= v[i];` means `sum = sum * v[i];` Other
`sum /= v[i];` means `sum = sum / v[i];`
`sum %= 2;` means `sum = sum % 2;`

operators can also be abbreviated this way.

Inline if

The following code:

```
if (r1 > r2) maxr = r1; else maxr = r2;
```

can be abbreviated:

```
maxr = (r1 > r2) ? r1 : r2;
```

Matrices - Details

For matrices allocated in the previous lecture

| | |
|---------------------------|---------------------------------------|
| <code>matrix</code> | points to a matrix |
| <code>matrix[0]</code> | points to the first row of the matrix |
| <code>matrix[0][0]</code> | The top left element of the matrix |

Which have the following types:

| | |
|---------------------------|------------------------|
| <code>matrix</code> | <code>double **</code> |
| <code>matrix[0]</code> | <code>double *</code> |
| <code>matrix[0][0]</code> | <code>double</code> |

Higher Dimensional Arrays

```
double *** alloc3Tensor(unsigned int dim)
{
    unsigned int i, j;
    double *** tensor;
    tensor = (double ***) malloc(dim*sizeof(double **));
    if (!tensor) return NULL;
    tensor[0] = (double **) malloc(dim*dim*sizeof(double **));
    if (!tensor[0])
    {
        free(tensor);
        return NULL;
    }
    tensor[0][0] = (double *) malloc(dim*dim*dim*sizeof(double));
    if (!tensor[0][0])
    {
        free(tensor[0]); free(tensor);
        return NULL;
    }
    for (i = 1; i < dim; i++)
        tensor[0][i] = tensor[0][i-1] + dim;
    for (i = 1; i < dim; i++)
    {
        tensor[i] = tensor[i-1] + dim;
        tensor[i][0] = tensor[i-1][0] + dim * dim;
        for (j = 1; j < dim; j++)
            tensor[i][j] = tensor[i][j-1] + dim;
    }
    return tensor;
}
```

More Multi-Index

- The 3D array behaves as expected:

```
tensor[i][j][k] = 0.5;
```

- The array can be freed with:

```
void free3Tensor(double *** tensor)
{
    free(tensor[0][0]);
    free(tensor[0]);
    free(tensor);
}
```

More Elaborate Data Structures

- In C we are able to manipulate data directly, this has allowed us to partition a contiguous array of memory into matrix rows.
- We needn't restrict ourselves to structures where the number of elements per row is constant, one example is Pascal's triangle:

```
      1
     1 2 1
    1 3 3 1
   1 4 6 4 1
      ⋮
```

```

int main()
{
    unsigned int size = 10, r, c;
    int ** pasc;
    pasc = (int **) malloc(size*sizeof(int));
    pasc[0] = (int *) malloc(size*(size+1)*sizeof(int)/2);
    /* check the mallocs */
    for (r=1; r < size; r++) pasc[r] = pasc[r-1]+r;
    pasc[0][0] = 1;
    for (r = 1; r < size; r++)
    {
        for (c = 1; c < r; c++)
            pasc[r][c] = pasc[r-1][c-1]+pasc[r-1][c];
        pasc[r][0] = 1; pasc[r][r] = 1;
    }
    for (r=0; r < size; r++)
    {
        for (c = 0; c < r+1; c++)
            printf("%3d ", pasc[r][c]);
        printf("\n");
    }
    free(pasc[0]);
    free(pasc);
    return 0;
}

```


Pascal's Triangle

- The program outputs:

```
1
1  1
1  2  1
1  3  3  1
1  4  6  4  1
1  5 10 10 5  1
1  6 15 20 15 6  1
1  7 21 35 35 21 7  1
1  8 28 56 70 56 28 8  1
1  9 36 84 126 126 84 36 9  1
```

- Whilst in memory, the data structure looks like:

| | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|-----|
| 1 | 1 | 1 | 1 | 2 | 1 | 1 | 3 | 3 | 1 | ... |
|---|---|---|---|---|---|---|---|---|---|-----|

Custom Data Types

Recall, when dealing with files we used:

```
FILE * file;  
file = fopen ( "myfile.txt", "r" );
```

FILE is in fact a custom data type, with its own size:

```
#include <stdio.h>  
  
int main()  
{  
    printf("sizeof(FILE) = %d\n", sizeof(FILE));  
    return 0;  
}
```

Except from MS - `<stdio.h>`

(From Visual Studio 2008)

```
struct _iobuf {  
    char *_ptr;  
    int  _cnt;  
    char *_base;  
    int  _flag;  
    int  _file;  
    int  _charbuf;  
    int  _bufsiz;  
    char *_tmpfname;  
};  
typedef struct _iobuf FILE;
```

typedef Structures : Custom Data Types

- FILE is a custom data type defined in `<stdio.h>`.
- It is comprised of elements of different, known, types.
- Each element of the FILE *structure* has its own name.

Let's consider a structure of our own, one of the simplest examples is a complex number:

```
typedef struct
{
    double real;
    double imag;
} complex;
```

C99

C99 fully supports its own complex type.

Definitions of structures take the following form:

```
typedef struct
{
    elementType elementName;
    elementType elementName;
    :
} structureTypeName ;
```

Handling Structures

- 1 A structure may be an argument to a function.
- 2 A function may return a structure.
- 3 A pointer may point to a structure.
- 4 Structures are referenced as normal: `&name`.
- 5 Elements of a structure are referenced as:
`&name.element`.
- 6 A pointer to a structure may be passed to a function.
- 7 If `p` is a pointer to a structure, then `p->member` allows us to access it's members.

Example Function Applying to Structures

```
#include <stdio.h>

typedef struct
{
    double real;
    double imag;
} complex;

void printComplex(complex * mc)
{
    printf("%lg + %lgi\n", mc->real, mc->imag);
}

int main()
{
    complex c1 = {1.0, 0.5}; /* assignment at declaration */
    printComplex(&c1);       /* pass pointer to struct */
    c1.real = 10.0;          /* piecewise assignment */
    printComplex(&c1);
    return 0;
}
```

More Structures

- Passing structures via pointer is usually more efficient.
- Assignment can be done at declaration or after.
- In C we are not allowed to overload operators, so the following won't work:

`c1 = c2 + c3;` (where `c1` and `c2` are `complex`)

so we need to do something like:

`complexAdd(&c1, &c2, &c3);`

- Arrays of `structs` are allowed (so matrices can consist of complex numbers for example).
- Structures can contain structures as elements.

Complex Number Support in C/C++

C++

Complex numbers are supported in C++ as `classes`, also the operators are overloaded properly too meaning any code which uses them will be concise. (look in `<complex>`).

C99

Complex numbers are supported as a native data type (not `struct`) in C99. Unfortunately not many compilers support this. GNU C, fully supports complex numbers (see `<complex.h>`).

C90

Complex number support in C90 is non-existent. I would recommend either third party libraries or a switch to C99/C++ for heavy complex number use.

Another Example struct - from <time.h>

```
struct tm
{
    int tm_sec;      /* seconds after the minute - [0,59] */
    int tm_min;      /* minutes after the hour - [0,59] */
    int tm_hour;      /* hours since midnight - [0,23] */
    int tm_mday;      /* day of the month - [1,31] */
    int tm_mon;       /* months since January - [0,11] */
    int tm_year;      /* years since 1900 */
    int tm_wday;      /* days since Sunday - [0,6] */
    int tm_yday;      /* days since January 1 - [0,365] */
    int tm_isdst;     /* daylight savings time flag */
};
```

- As seen in the first lecture, C programs are *compiled* into a low level (machine specific language).
- This language is called *assembly language*.
- On PCs and Macs Intel x86 assembler is ubiquitous.
- Most C compilers allow you to view the assembler that they generate.

addMatrices revisited - C code

From the last lecture we saw a function to add two matrices together:

```
void addMatrices(double ** matrixA, double ** matrixB,
                 double ** matrixR,
                 unsigned int rows, unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++)
        for (j = 0; j < cols; j++)
            matrixR[i][j] = matrixA[i][j]+matrixB[i][j];
}
```

When we compile this, we get...

addMatrices revisited - when compiled...

```
_addMatrices PROC
; 35 : {
        push    ebp
        mov     ebp, esp
        sub     esp, 216
        push    ebx
        push    esi
        push    edi
        lea     edi, DWORD PTR [ebp-216]
        mov     ecx, 54
        mov     eax, -858993460
        rep stosd
; 36 :     unsigned int i, j;
; 37 :     for (i = 0; i < rows; i++)
        mov     DWORD PTR _i$[ebp], 0
        jmp     SHORT $LN6@addMatrice
$LN5@addMatrice:
        mov     eax, DWORD PTR _i$[ebp]
        add     eax, 1
        mov     DWORD PTR _i$[ebp], eax
$LN6@addMatrice:
        mov     eax, DWORD PTR _i$[ebp]
        cmp     eax, DWORD PTR _rows$[ebp]
        jae     SHORT $LN4@addMatrice
; for (j = 0; j < rows; j++)
        mov     DWORD PTR _j$[ebp], 0
        jmp     SHORT $LN3@addMatrice
$LN2@addMatrice:
        mov     eax, DWORD PTR _j$[ebp]
        add     eax, 1
        mov     DWORD PTR _j$[ebp], eax

$LN3@addMatrice:
        mov     eax, DWORD PTR _j$[ebp]
        cmp     eax, DWORD PTR _rows$[ebp]
        jae     SHORT $LN1@addMatrice
;matrixR[i][j] = matrixA[i][j]+matrixB[i][j];
        mov     eax, DWORD PTR _i$[ebp]
        mov     ecx, DWORD PTR _matrixA$[ebp]
        mov     edx, DWORD PTR [ecx+eax*4]
        mov     eax, DWORD PTR _i$[ebp]
        mov     ecx, DWORD PTR _matrixB$[ebp]
        mov     eax, DWORD PTR [ecx+eax*4]
        mov     ecx, DWORD PTR _j$[ebp]
        mov     esi, DWORD PTR _j$[ebp]
        fld     QWORD PTR [edx+ecx*8]
        fadd    QWORD PTR [eax+esi*8]
        mov     edx, DWORD PTR _i$[ebp]
        mov     eax, DWORD PTR _matrixR$[ebp]
        mov     ecx, DWORD PTR [eax+edx*4]
        mov     edx, DWORD PTR _j$[ebp]
        fstp    QWORD PTR [ecx+edx*8]
        jmp     SHORT $LN2@addMatrice
$LN1@addMatrice:
        jmp     SHORT $LN5@addMatrice
$LN4@addMatrice:
        pop     edi
        pop     esi
        pop     ebx
        mov     esp, ebp
        pop     ebp
        ret     0
_addMatrices ENDP
```

- A few lines of C becomes > 30 lines of assembler (only three of which are actually floating point instructions!).
- It is possible to write a much smaller assembler routine by hand ≈ 10 -20 instructions long.
- This would run ≈ 3 times quicker than the C compiled routine (this is a general rule of thumb).
- Any custom assembly code would only target a very specific chip, however.

- Rather than rewrite the assembly code, it is easier to ask the C compiler to perform code optimisation itself.
- By default C will compile the code as it appears (the exact order of operations is preserved etc), this aids debugging.
- C compilers can be told to optimise their code in the following ways:
 - MSVC Project configuration options can be set, defaults in “Release” build do a good job.
 - gcc The `-O` command line flags influence optimisation, `-O0` means “off” whilst `-O3` means “extremely aggressive”.

Optimisation Example - Visual Studio 2008

Debug Build

```
$LN2@addMatrice:
    mov     eax, DWORD PTR _j$[ebp]
    add     eax, 1
    mov     DWORD PTR _j$[ebp], eax
$LN3@addMatrice:
    mov     eax, DWORD PTR _j$[ebp]
    cmp     eax, DWORD PTR _rows$[ebp]
    jae     SHORT $LN1@addMatrice
;matrixR[i][j] = matrixA[i][j]+matrixB[i][j];
    mov     eax, DWORD PTR _i$[ebp]
    mov     ecx, DWORD PTR _matrixA$[ebp]
    mov     edx, DWORD PTR [ecx+eax*4]
    mov     eax, DWORD PTR _i$[ebp]
    mov     ecx, DWORD PTR _matrixB$[ebp]
    mov     eax, DWORD PTR [ecx+eax*4]
    mov     ecx, DWORD PTR _j$[ebp]
    mov     esi, DWORD PTR _j$[ebp]
    fld     QWORD PTR [edx+ecx*8]
    fadd    QWORD PTR [eax+esi*8]
    mov     edx, DWORD PTR _i$[ebp]
    mov     eax, DWORD PTR _matrixR$[ebp]
    mov     ecx, DWORD PTR [eax+edx*4]
    mov     edx, DWORD PTR _j$[ebp]
    fstp    QWORD PTR [ecx+edx*8]
    jmp     SHORT $LN2@addMatrice
```

Release Build

```
;matrixR[i][j] = matrixA[i][j]+matrixB[i][j];
    fld     QWORD PTR [edx+eax]
    fadd    QWORD PTR [eax]
    fstp    QWORD PTR [esi+eax]
    add     eax, 8
    dec     ebx
    jne     SHORT $LL3@addMatrice
```


Bits and Bytes

Bytes

Smallest *addressable* unit of memory, each byte is composed of eight bits.

Bits

These are the smallest units of computer memory, each bit can be either 0 or 1.

- Addressing bits individually requires some extra operations to be carried out.
- There are good reasons for accessing data at the bit-level however.

Efficient Data Packing

Given a 32 million base pair chromosome

It will require: ≈ 64 megabytes to store as `short`
 ≈ 32 megabytes to store as `char` (next lecture)
 ≈ 8 megabytes to store as bit data.
(i.e. a 2-4 year research lead if following Moore's Law).

Computer Graphics

Given a monochrome print image of 2400 dpi rendered over 80 square inches gives $\approx 500,000,000$ dots. This takes up:

≈ 1 gigabyte if using `short`
 ≈ 64 megabytes if using bits.

Bit Manipulation Friendly Data Types

As seen before the unsigned integer data types have values ranging from 0 to $2^n - 1$ where n is the number of bits in the data type. Some examples (for my machine):

| Data Type | <i>n</i> |
|------------------------|----------|
| unsigned short | 16 |
| unsigned int | 32 |
| unsigned long int | 64 |
| unsigned long long int | 128 |

Unsigned data types are also desirable for accessing array indices as they can never be negative.

How to Get Them In and Out of The Computer

- They can be read using `scanf` and `"%u"`, `"%lu"` or `"%Lu"`.
- They can be printed using `printf` and `"%u"` or `"%Lu"`.
- We can output to octal (base 8 numbers) using `printf` and `"%o"`
- Also we can output to hexadecimal (base 16, 0-9 and a-f) using `printf` and `"%x"` or `"%X"`.

Example - byte (char)

| | |
|-------------|----------|
| Binary | 10101011 |
| Hexadecimal | AB |
| Decimal | 171 |
| Octal | 253 |

Manipulating Bits within an Unsigned Integer

- C can shift all the bits comprising a number a fixed number of places to the left or right.
- Zeros are propagated in to the vacated spaces.
- Bits that shift outside, disappear. (i.e. the shift is not cyclic).
- Bit shifting is accomplished with the `>>` (right) and `<<` (left) operators.

For example:

$$1 \ll 2 = 4$$

$$8 \gg 3 = 1$$

Bit shifts are much cheaper than multiplying or dividing by powers of two.

4 Bitwise operators $\&$, $|$, \wedge and \sim

Assuming 0 is false and 1 is true, we have the following bitwise logical operators.

And Operator($\&$)

| | | | | |
|---------|---|---|---|---|
| N1 | 0 | 0 | 1 | 1 |
| N2 | 0 | 1 | 0 | 1 |
| N1 & N2 | 0 | 0 | 0 | 1 |

Or Operator($|$)

| | | | | |
|---------|---|---|---|---|
| N1 | 0 | 0 | 1 | 1 |
| N2 | 0 | 1 | 0 | 1 |
| N1 N2 | 0 | 1 | 1 | 1 |

Exclusive Or(\wedge)

| | | | | |
|----------------|---|---|---|---|
| N1 | 0 | 0 | 1 | 1 |
| N2 | 0 | 1 | 0 | 1 |
| N1 \wedge N2 | 0 | 1 | 1 | 0 |

Not Operator(\sim)

| | | |
|-----------|---|---|
| N1 | 0 | 1 |
| \sim N1 | 1 | 0 |

Case Study: The Sieve of Eratosthenes

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 5 | 7 | 9 | 11 | 13 | 15 | 17 | 19 | 21 | 23 | 25 | 27 | 29 | 31 | 33 | 35 | 37 | 39 | 41 | 43 | 45 | 47 | 49 | 51 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|

This gives the primes (we add 2 to the beginning of the list):

2 3 5 7 11 13 17 19 23 29 31 37 41 43 47

The Sieve of Eratosthenes: In C

- When implementing this in C it makes sense to use a bit to indicate 'primeness' of a number.
- The smallest addressable unit of memory in C is the `char` which consists of 8 bits.
- We therefore need to perform *masking* to isolate individual bits.

Masking

We access the j^{th} bit of a variable `x` as follows:

| | |
|--|--------------------------|
| <code>if (x & (1 << j))</code> | Check to see if it's set |
| <code>x = (1 << j)</code> | To set the bit |
| <code>x &= ~(1 << j)</code> | To clear the bit |

The Sieve of Eratosthenes: Implementation

```
void findPrimes(NUM_TYPE * Prime_List, int max_num)
{
    int current_num = 3;
    NUM_TYPE Mask=1;

    while(current_num*current_num <= max_num)
    {
        int Pnum = current_num/(2 * BITS_PER_NUM);
        int Pbit = (current_num-Pnum * 2 * BITS_PER_NUM)/2;
        /* Current Number is prime so strike out multiples */
        if(~Prime_List[Pnum] & (Mask <<Pbit))
        {
            int strike = current_num * current_num;
            while(strike <= max_num)
            {
                int Snum = strike/(2*BITS_PER_NUM);
                int Sbit = (strike - Snum * 2 * BITS_PER_NUM)/2;
                Prime_List[Snum] |= (Mask << Sbit);
                strike += 2 * current_num;
            }
        }
        current_num += 2;
    }
}
```

The Sieve of Eratosthenes : Printing out the Primes

```
void printPrimes(NUM_TYPE * Prime_List, int max_num)
{
    int i;
    NUM_TYPE Mask=1;

    for(i = 3;i <= max_num;i = i+2)
    {
        int Pnum=i/(2*BITS_PER_NUM);
        int Pbit=(i-Pnum*2*BITS_PER_NUM)/2;
        if (~Prime_List[Pnum] & (Mask <<Pbit))
            printf("%d\n", i);
    }
}
```